

Empowering Meta-Analysis: Leveraging Large Language Models for Scientific Synthesis

Jawad Ibn Ahad Apurba-NSU R&D Lab, ECE North South University Dhaka, Bangladesh jawad.ibn@northsouth.edu	Rafeed Mohammad Sultan Apurba-NSU R&D Lab, ECE North South University Dhaka, Bangladesh rafeed.sultan@northsouth.edu	Abraham Kaikobad Apurba-NSU R&D Lab, ECE North South University Dhaka, Bangladesh abraham.kaikobad@northsouth.edu	Fuad Rahman Apurba Technologies Sunnyvale, CA 94085, USA fuad@apurbatech.com
Mohammad Ruhul Amin Computer and Information Science Fordham University New York, USA mamin17@fordham.edu	Nabeel Mohammed Apurba-NSU R&D Lab, ECE North South University Dhaka, Bangladesh nabeel.mohammed@northsouth.edu	Shafin Rahman Apurba-NSU R&D Lab, ECE North South University Dhaka, Bangladesh shafin.rahman@northsouth.edu	

Abstract—This additional material offers supplementary details that bolster the conclusions outlined in the main manuscript.

- **Section 1: Experimental Details** (additional discussion in support of Section 4.A and 3.C of the main paper).
- **Section 2: Evaluation** (additional discussion in support of Section 4.A of the main paper).

I. EXPERIMENTAL DETAILS:

Table I provides a detailed breakdown of the hyperparameters we used in our experiments. The two models underwent ten epochs of fine-tuning, with a batch size of four and a learning rate of $2e-4$. The Supervised Fine-tuning Trainer (SFTTrainer) with customized loss metrics, led the entire training process.

A. Set up for RAG

Setup for RAG: For the retrieval augmented generation approach, sentence transformers and embeddings handle large context lengths. The data is processed using the "Recursive TextSplitter" from *LangChain*¹ to break the text into manageable chunks, satisfying LLMs maximum context length. Subsequently, the *HuggingFaceEmbeddings* model "sentence-transformers/all-MiniLM-L6-v2" generates embeddings for the text chunks. Sentences are tokenized and converted into numerical tokens using the *AutoTokenizer* from *HuggingFace*². The pre-trained Sentence Transformer model "all-MiniLM-L6-v2" obtains encoded inputs, including padding and truncation, as specified in the tokenizer. The normalized embeddings are stored in *Pinecone*³, creating a searchable index. A question-answering chain, developed as a *RetrievalQA*, retrieves information based on the query. For the RAG approach, the dataset includes a third column containing specific queries regarding

the context, effectively framing the query as a question about the contents of the meta-papers' abstracts.

Vector database: For the RAG approach, a vector database is necessary to store the embeddings of provided large contexts. The normalized embeddings are stored in *Pinecone*³, creating a searchable index. *Pinecone* serves as the vector knowledge base, storing the embeddings from the sentence transformers.

Question Answering Chain: A question-answering chain, developed as a *RetrievalQA*, retrieves information based on the query. For the RAG approach, the dataset includes a third column containing specific queries regarding the context, effectively framing the query as a question about the contents of meta-papers' abstracts.

B. Set up for combining RAG with Fine-tuned LLM

Combining Retrieval Augmentation (RAG) with fine-tuned models enhances Language Model Models' (LLMs) performance, especially for tasks like generating meta-analysis abstracts. Fine-tuning LLMs improves task-specific capabilities but struggles with large contexts. RAG acts as an information bridge, retrieving relevant knowledge from external sources, enhancing accuracy, reducing fine-tuning needs, and mitigating irrelevant content generation. RAG ensures semantic alignment between retrieved information and user queries, improving the quality, readability, and clarity of meta-analysis abstracts. This integrated approach produces more accurate, focused, and readable content.

Furthermore, while fine-tuning large language models (LLMs) allows them to excel at specific tasks like generating meta-analysis abstracts, their inherent limitation in handling large contexts becomes a hurdle. This is where combining retrieval augmentation (RAG) with fine-tuning offers a powerful solution. RAG acts like an information bridge, retrieving relevant knowledge from external sources beyond the limited chunks provided. This expanded knowledge base empowers the LLM to process complex information from various sources, leading

¹LangChain: <https://www.langchain.com/>

²HuggingFace: <https://huggingface.co/>

³Pinecone: <https://www.pinecone.io/>

TABLE I: Hyperparameter values used for our experiment

Hyperparameter	Value	Hyperparameter	Value
Batch Size	4	Learning Rate	2e-4
Epochs	10	Temperature	0.7
Loss Function	ICD	Optimizer	AdamW
Lora Alpha	4	Lora Dropout	0.1
Rank (r)	64	–	–

to several advantages. Firstly, the accuracy and factual grounding of the generated abstracts are significantly enhanced. Since factual accuracy is crucial for meta-analysis, RAG ensures the LLM incorporates reliable external knowledge. Secondly, leveraging external information through RAG has the potential to reduce the extensive fine-tuning required for LLMs, leading to faster training times. Finally, RAG helps mitigate the issue of irrelevant content generation that can sometimes plague fine-tuned models. In Retrieval-Augmented Generation (RAG) systems, semantic search ensures retrieved information is not just relevant, but also semantically aligns with the task. RAG uses semantic similarity to identify contexts that closely match the user’s query and the expertise of the fine-tuned large language model (LLM). This injects relevant information, enhancing the LLM’s understanding and synthesis of meta-analysis abstracts, leading to higher-quality outputs. Ultimately, RAG contributes to improved readability and clarity in the final abstracts. This combined approach, where fine-tuning grants the LLM the ability to extract relevant patterns for meta-analysis generation and RAG broadens the context with the necessary information, leads to more accurate, focused, and readable meta-analysis content.

This integrated process can reduce hallucinations and produce more accurate answers efficiently. The fine-tuned models are employed to perform RAG on unseen data. As aforementioned above, Figure 2 illustrates the complete workflow for this approach.

C. Data Processing and Context Analysis

The dataset, comprising main papers and their associated metadata, underwent preprocessing and context analysis which is shown in Fig1. The preprocessing steps involved chunking the metadata using a custom text splitter, designed to partition texts into manageable segments. Each segment was associated with its corresponding main paper label.

Subsequently, the length of each chunk, measured in characters, was computed for both the main paper and its metadata.

As previously mentioned, it will not be feasible to send all the information at once because the context size for abstracts will be greater than the context length of LLMs. Figure 2 illustrates the possibility of providing input from a single chunk for only fine-tuning inference. Thus, it goes without saying that the model will be unable to extract any data from the other chunked contexts. On the fine-tuned model, RAG is applied to solve this issue. This will make it possible to insert the entire context, including each abstract. This

extensive context will be subjected to a semantic search process using queries, and the context that is found will be sent as input to LLM. LLM will acquire knowledge about each chunk in this manner. It is observed that through the fine-tuning process, models learned the patterns of constructing meta-analysis abstracts. Additionally, by applying RAG, they were able to learn from large unseen contexts, resulting in better performance.

II. EVALUATION

Human Evaluation. In this section, we outline the evaluation process employed in this study. Following the generation of responses by LLMs, a human evaluation process is conducted to align the generated text with human judgment. This process, adapted from a previous study, involves human judges categorizing LLM-generated text into three categories: "Relevant" (2), indicating a close resemblance to the ground truth with high similarity and inclusion of important input information; "Somewhat-Relevant" (1), suggesting an acceptable similarity with valuable information within an acceptable margin; and "Irrelevant" (0), signifying a lack of important information or the presence of irrelevant contexts. This classification framework ensures a rigorous assessment of generated meta-analysis abstracts against expected standards.

The evaluation process involves three individuals independently assessing statements and responses from each model output as shown in Fig3, which was done in Google Sheets. These evaluators, instructed to evaluate relevancy, conduct their assessments without access to each other’s evaluations. Final results are determined through hard voting by a designated final evaluator. If two of the evaluators select "Relevant" (2), the final vote is considered as 2; the same applies for "Irrelevant" (0). However, in cases where there are different evaluation labels, such as 0, 1, and 2, and there is no agreement among the evaluators, the final decision is assigned as 1. This is because "Somewhat-Relevant" (1) is considered almost similar to "Relevant" (2), indicating some level of relevancy in the processed output of models. Thus, if one evaluator labels it as 2 (relevant) and another as 1 (somewhat-relevant), the processed output is deemed somewhat-relevant, resulting in a final vote of 1. To mitigate biases, the human evaluation is conducted by university students who are not authors of the study.

Evaluation with RAG

Fig. 4 displays a sample for evaluation for the Retrieval Augmented Generation (RAG) approach. Cosine-similarity

```

import pandas as pd
from langchain_text_splitters import RecursiveCharacterTextSplitter

# Instantiate RecursiveCharacterTextSplitter with custom parameters
text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=2000,
    chunk_overlap=200,
    length_function=len,
    #separators=['\n\n','MP:', '\n', ' ', '']
)

# Function to split text into chunks using the text splitter
def split_text_into_chunks(text):
    return text_splitter.split_text(text)

chunked_df = pd.DataFrame({"Paper": df["Main_Paper"], "Meta_Paper": df["MP"]})
chunked_data = {"Paper": [], "Meta_Paper": []}

# Iterate through each row in the dataset
for index, row in df.iterrows():
    context = row["MP"]
    label = row["Main_Paper"]

    # Split the context into overlapping chunks
    chunks = split_text_into_chunks(context)
    for chunk in chunks:
        chunked_data["Paper"].append(label)
        chunked_data["Meta_Paper"].append(chunk)

# Create a new DataFrame with the chunked data
chunked_df = pd.DataFrame(chunked_data)

```

Fig. 1: Illustration of the Chunking Process. The text splitter algorithm segments the metadata into overlapping chunks, facilitating efficient processing and analysis of the dataset

is used to evaluate the produced meta-analysis using RAG because the loss function for fine-tuning LLM was designed using dissimilarity. Again, due to the approach of RAG, LLMs are bound to produce contexts that are expected to be relevant to the input context. So, the generated outputs via RAG are relevant meta-analysis abstracts. For this reason, we need to calculate the similarity between two abstracts: one is an actual meta-analysis abstract, and the other is a generated abstract in the RAG approach which is the processed output. The average

cosine similarity is listed in the paper.

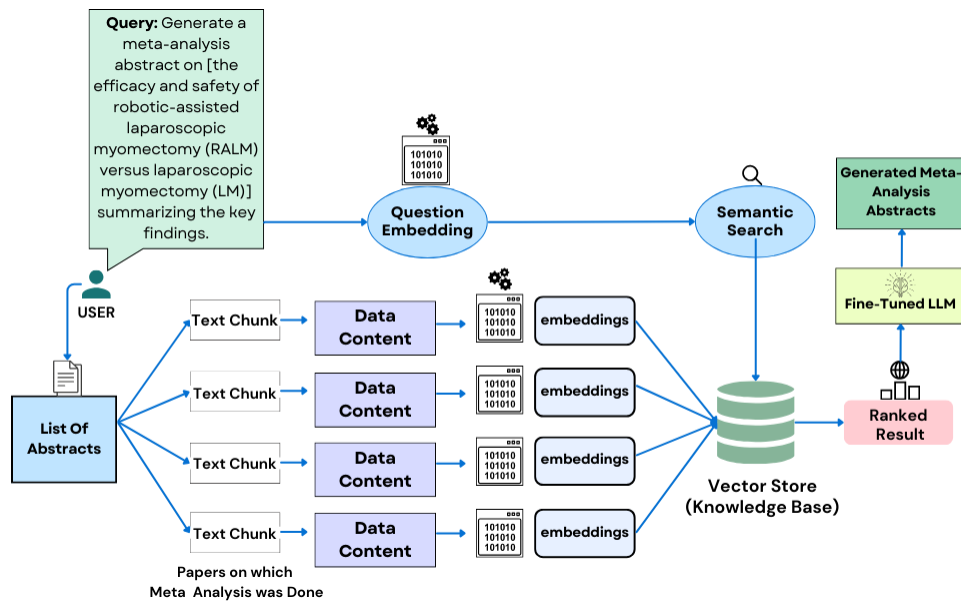


Fig. 2: The user will submit a query along with abstracts from several publications and a broad context. There will be several segments that are taken out of context. Every chunk's vector embeddings will be kept in a vector database. Using the embedded query in the vector base, a semantic search will be carried out. The context that is retrieved will be fed to the fine-tuned LLM that was trained on the conventions for generating meta-analysis abstracts. Lastly, the model produces an abstract for a meta-analysis.

	A ▼	B	C	D	E	F
1	Meta_Paper	Processed Output	Evaluation-1	Evaluation - 2	Evaluation - 3	Final_Decision
2	MP: Assessment	The present met	2	2	2	2
3	of Coronavirus d	The provided ab Some key finding * Improved over * Reduced durat * Accelerated re * Mortality redu Some numerical * Overall surviva * Duration of ho * CRS resolution * Mortality rate: * Hazard ratio fo Please note that	2	2	2	2
4	days of illness w	A meta-analysis	1	2	1	1
5	MP: BACKGROU	The provided ab	2	2	2	2
6	(95%-CI 0.66–1.0	Based on the pro	2	2	1	2
7	significant differ	The provided ab	2	2	2	2
8	(MRS) was used	Based on the pro Some numerical	2	2	2	2
9	model to predict	The present met	2	2	2	2

Fig. 3: Human evaluation spreadsheet showing meta-analysis abstract papers, LLM-generated outputs, evaluations of three individuals, and the overall evaluation.

	A	B	C	D
1	Ground Truth	MP	Processed Output	sim
2	Background The available evidence about the possible	MP: ObjectiveSeveral random	This meta-analysis abstract summarizes the findings of several	0.7998
3	Introduction Several studies have demonstrated that	MP:Background: Depression	In recent studies, interventions have been explored for improving	0.7519
4			The provided abstracts focus on the impact of physical activity The results of the meta-analysis showed that physical activity Moreover, the meta-analysis results indicated that the impact	
	Background Glycemic control is vital to patient care, and	MP:Purpose Diabetes mellitus	Overall, the results of the meta-analysis suggest that engaging	0.8048
5			Meta-analysis abstract:	
	Objective This meta-analysis aimed to evaluate the effect	MP: Background Recent in	Pioglitazone use is associated with a reduced risk of Parkinson's	0.9038
6	Background Recurrent hemarthrosis is one of the major	MP:Background: The primary	Myofascial therapy has been investigated as an effective treatment	0.8022
7	Background Curcumin, a natural polyphenolic compound	MP:Polycystic ovary syndrome	A systematic review and meta-analysis of randomized controlled	0.8189
8	Purpose Despite the high prevalence of postdialysis fatigue	MP:Background The aetiology	This meta-analysis abstract summarizes the key findings of several	0.7565
9	Purpose Research has shown that prone positioning (PP)	MP:OBJECTIVES: To determine	Prone positioning (PP) was found to reduce mortality in patients	0.8481
10	Handhelds (e.g., cell phones, tablets) are promising for	MP:The importance of mobile	Mobile technologies have become increasingly important in education	0.873
11	Background Circulating concentrations of homocysteine	MP:Objectives To investigate	A meta-analysis was conducted to investigate the relationship between	0.805
12			Meta-analysis Abstract:	
	Introduction Previous randomized controlled trials (RCTs)	MP: Background Vitamin D	A systematic review and meta-analysis of randomized clinical	0.8601
			Functional neuroimaging and Dispositional Variables for Clinical	
			=====	

Fig. 4: Evaluation of processed output using RAG on fine-tuned Mistral-v0.1 (7B) LLM is shown here. The last column 'sim' means similarity with ground truth (SWGT), which refers to cosine-similarity between the generated meta-analysis abstract and the actual meta-analysis abstract.